Daily Proof of Liabilities

Antoine Cyr

A Thesis in

The Concordia Institute for Computer Science (MCompSc)

Presented in Partial Fulfillment of the Requirements

For the Degree of

Master

(Computer Science)

at

Concordia University

Montréal, Québec, Canada

September 2024

CONCORDIA UNIVERSITY
School of Graduate Studies

This is to certify that the thesis prepared

By: **Antoine Cyr**

Entitled: **Title**

and submitted in partial fulfillment of the requirements for the degree of

Master (Computer Science)

complies with the regulations of this University and meets the accepted standards with respect to originality and quality.

Signed by the final examining committee:

———————————————————————— Chair

———————————————————————— External Examiner

———————————————————————— Examiner

———————————————————————— Examiner

———————————————————————— Examiner

———————————————————————— Supervisor

Approved by ————————————————————————————

01 Sept 2024 ————————————————————————————

Abstract

Name: **Antoine Cyr**

Title: **Daily Proof of Liabilities**

A proof of solvency's goal is to demonstrate that a cryptocurrency exchange possesses sufficient funds to satisfy client withdrawals. In this thesis, we introduce an improvement to the prevailing way of building a proof of liabilities. We use the Nova novel way of proving that a balance is included in the proof of liabilities (i.e. proof of inclusion), and apply it to the proof of liabilities itself. We use the circuit designed to show the proof of inclusion of a Merkle tree, and modify it to prove a list of balance changes in the Merkle tree. While this is slower than producing the whole Merkle tree when you have many changes, this new circuit design enables to separate the proof into multiple smaller proofs, enabling the use of the Nova folding scheme. The folding of arithmetics circuits reduces the computation needed for a daily proof of liabilities, enabling the possibility of obtaining this proof at a higher frequency, potentially as frequently as every block.

Acknowledgments

Contents

List of Figures

Chapter 1

Introduction

In the context of cryptocurrencies, the trust between marketplaces and users is at an all time low. Following the recent bankruptcy and mishandling of customer funds by marketplaces like FTX, users want and need to know that marketplaces have all the funds in their possession. Traditional financial institutions commonly rely on audits as the established method to demonstrate their solvency. The outcomes of third-party examinations gain universal acceptance due to the trust placed in these third parties. However, auditing a cryptocurrency marketplace presents certain challenges. Audits are not a scalable solution because the funds can be moved around more quickly than in the traditional finance world. This would call for a high-frequency auditing approach, perhaps even daily audits. Here, automating the process becomes advantageous.

Another challenge pertains to trust in the third party, which must be mutual. The institution needs to trust that their data won't be compromised, and the public must

trust the authenticity of the audit results. This is where zero-knowledge proofs come into play.

The proof of solvency uses zero-knowledge to demonstrate, without revealing any information, that the assets in control are greater than the liabilities. The proof of solvency is made up of a proof of asset and a proof of liabilities, a structure that is self-explanatory.

This paper extends the work done on proof of liabilities only. The most prevalent way to implement a proof of liabilities is by using a Merkle Sum tree. In a Merkle tree, every parent node represents the hash of its child nodes. In the Merkle Sum tree, each node additionally holds the balance information. The leaf nodes consist of user IDs along with their respective balances, while the root node contains the hash of the entire tree, summarizing the total balance.

Alongside the proof of liabilities, there is a proof of inclusion, where you demonstrate that for each customer, their balance is included in the tree. This is achieved by proving the Merkle path, which is sufficient for validation.

This paper explores a way to minimize the cost of doing these proofs everyday, while maintaining complete privacy (i.e the Merkle tree is kept private).

1.1   Outline

Chapter 2 gives the background information needed to understand this paper. It elaborates on Bitcoin, including its transactions, the network, proof-of-work, and Merkle

tree. Marketplaces are also discussed. Additionally, the chapter delves into Zero Knowledge, covering non-interactive proofs, SNARKs, and arithmetic circuits. The concept of proof of solvency, with a focus on real-world applications, is introduced.

Chapter 3 goes deeper into zero knowledge, more specifically the different recursion techniques, which will be used to reduce the proof size and make daily proofs easier. The techniques include aggregation, recursion schemes, accumulation and folding scheme.

Finally, Chapter 4 explores the proof and circuit construction of the thesis. We start with a simple circuit for the proof of liabilities, as well as a circuit for the proof of inclusion for a Merkle tree. Next, we modify the initial proof of inclusion to prove the inclusion of many Merkle trees at the same time, using the folding scheme. The folding scheme is a recent technique brought to us by Nova that allows to "fold" many circuits together, reducing the proof size. This new inclusion circuit was implemented by them using their technique.

For the proof of liabilities, we adopt an alternate strategy aiming to reduce the size of generating the proof on the second day, given the original Merkle tree is already established. This is the first novel proof of the thesis. This involves developing a secondary liabilities circuit that constructs a Merkle tree, utilizing an initial Merkle tree alongside the changes as input. We then optimize it using the folding scheme.

1.2 Motivation

The novel aspect of this thesis emerged after noticing the significant boost that the Nova folding scheme provided to the proof of inclusion. This observation led me to ask: how can this be applied elsewhere? The answer is straightforward: the key component of the proof of liabilities, which is the proof of liabilities itself, can benefit from this approach.

Given how obvious this seems, why hasn't anyone else done it? There are two main reasons. First, the traditional method of conducting a proof of liabilities doesn't benefit from the Nova folding scheme. This thesis introduces a new circuit specifically designed to be used with the folding scheme. Without it, the circuit wouldn't be scalable, as the proofs would grow with each change. Second, the folding scheme is relatively new, and researchers are still exploring its full potential and all its applications.

1.3 Results

The proof of liabilities circuits are evaluated and compared based on the number of constraints and proof time. In this thesis, we demonstrate that the new liabilities circuit performs better when the number of changes is 0.05% of users or less. The new circuit is also better the fewer users we have, but it performs well enough with a high number of users. For instance, at 1 million users, the new circuit is 3 times smaller than the old circuit (1 000M constraints vs 3 000M constraints). At 1000

users, the results are way better; we have 300k constraints vs 3M constraints.

Additionally, we show that when combined with the Nova folding scheme, the new circuit remains superior even with 1% of changes. The folding circuit performs better the more users we have. For instance at 256 users the proof time is similar(around 50 seconds), but at 1M users we have around 20 000 seconds vs 200 000 seconds.

The percentage of changes is defined by the number of root changes in the tree between a previously proven Merkle tree and the targeted Merkle tree.

Chapter 2

Background

In the dynamic landscape of cryptocurrencies, ensuring transparency, accountability, and financial stability is paramount for marketplaces. This section provides an overview of the concepts and mechanisms necessary to follow along the conception of a daily proof of liabilities. We explore the various approaches adopted by leading exchanges to demonstrate their financial health through proof of reserves or solvency mechanisms. Additionally, we highlight the shortcomings and challenges associated with current solvency verification practices, mainly the lack of recurrent proof of reserves, paving the way for a daily proof in the subsequent sections.

2.1 Bitcoin

Bitcoin is recognized as the world's first successful cryptocurrency and decentralized digital currency. The goal of Bitcoin is to allow financial transactions to be settled on

their own, without the need of a middleman, typically financial institutions. Bitcoin is built on a peer-to-peer network, which means that every participant helps in safe keeping the transactions history, and propagating the new transactions. There is no single point of failure. This allows for transactions to occur in real time, in contrast to the delays encountered in the traditional finance world. Bitcoin defines two different concepts: Bitcoin the cryptocurrency, and Bitcoin the blockchain. The cryptocurrency resides on the blockchain. The Bitcoin blockchain is a decentralized ledger that records all Bitcoin (the cryptocurrency) transactions immutably and transparently. This blockchain serves as a verifiable record of all Bitcoin transactions, accessible to every participant in the network. The transparency afforded by the public blockchain engenders trust and accountability. [5]

2.1.1 Transactions

For every participant of the network, there is a public key, a private key and a wallet address associated with the participant. The public key is derived from the private key using elliptic curve multiplication, and the wallet address is derived from the public key using a hashing function. Both are one way functions, meaning you cannot derive the other way around. The public key serves as the unique identifier in the network, but it is the wallet address that typically defines a participant. The wallet address can be seen as a bank account number. When you send Bitcoin to someone, you send it to their wallet address. To be able to send some Bitcoin, you need to create a transaction and send it to the network. When transactions are sent on the

network, there is no way of knowing who propagated the transaction first. We need to make sure a transaction originates from the sender. The way to do that is to sign your transaction. The digital signature is created from the transaction data and the private key, which is only known by the owner of the address. The digital signature is created using the ECDSA (Elliptic Curve Digital Signature Algorithm) over the secp256k1 curve, which is the specific elliptic curve used by Bitcoin. We verify a Bitcoin transaction's digital signature by using the sender's public key to check that the signature matches the hashed transaction data. This is done through ECDSA, where the signature is compared to a computed value derived from the transaction hash and the public key. If they match, it confirms the transaction was signed with the corresponding private key, validating the transaction. Sending a transaction is the easiest problem to solve. The real challenge is to keep track of who owns what, and to avoid the double spending problem. The methodology for managing this is to keep the history of every single transaction. The transactions are bundled up into blocks, and the chain of blocks creates the blockchain. [5]

2.1.2 Network

The challenge of the network is to have every single node achieve consensus on the transaction history. Nodes are computers connected to the network, working on publishing new blocks. The nodes work collectively to establish order of transactions (sequencing). Every new transaction is broadcast to all nodes. The nodes put the transactions into a block, and try to publish that block. In order to publish a block,

each node needs to solve a proof-of-work challenge. When a node solves the challenge, it broadcasts the block to every node. The node accepts the block if all transactions are valid. There is no formal way of approving a new block. A node shows its acceptance by starting to work on a new block using the hash of the accepted block as previous hash. Some nodes might accept different blocks, if multiple blocks are propagated at the same time. This would create multiple chains. To solve this issue, the longest chain is considered to be the correct one. If two chains have the same length, nodes keep working on their respective chains until one of the chains receives a new block, breaking the tie. [5]
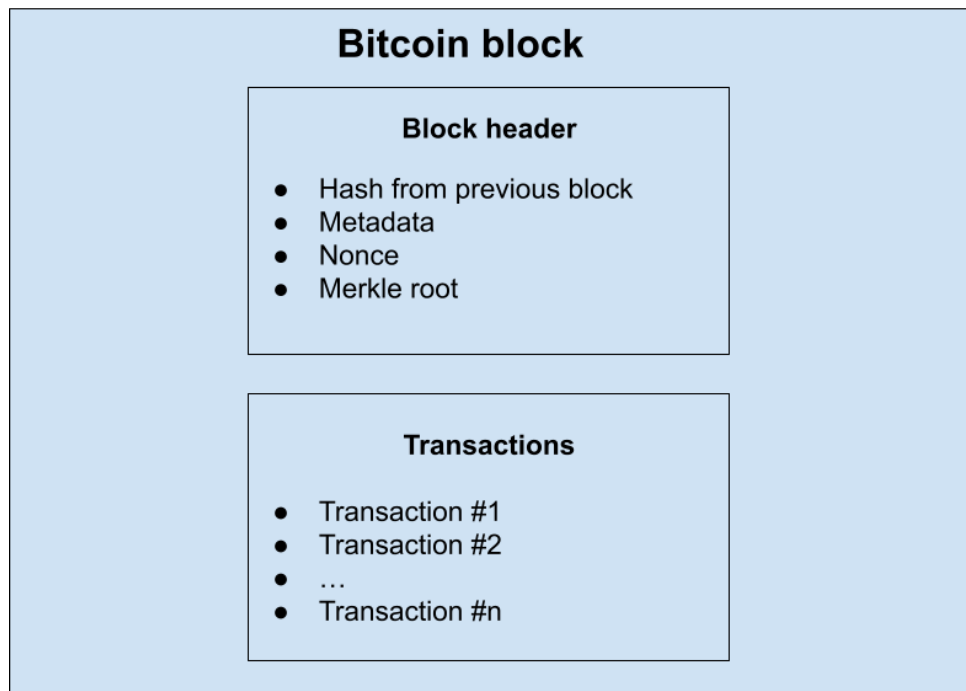

Bitcoin block

Block header

- Hash from previous block
- Metadata
- Nonce
- Merkle root

Transactions

- Transaction #1
- Transaction #2
- …
- Transaction #n

Figure 2.1: Bitcoin block

2.1.3 Proof-of-work

In order to submit a new block, a node has to find a hash with a specific number of leading zeros. The hash is produced by the SHA-256 algorithm and is 256 bits long, typically represented as a 64-character hexadecimal string. The number of leading zeros required is determined by the network's difficulty, which sets the target value that the hash must be less than or equal to for the block to be valid. For instance, the threshold might be $0000000000000000abc123...$ This cryptographic puzzle serves as a barrier to entry, ensuring that a lot of computational power was spent on creating the block. The way to test different hash values is to change the block timestamp, and the nonce value. The nonce value is there solely for that purpose. Once a block is published, you cannot change any value inside of it because the hash value would change. The immutability of the older blocks is what makes Bitcoin secure. To modify a block in the middle of the chain, you would need to redo the work of every single block made after that. The longest chain is determined by the cumulative proof-of-work invested in it. This is why we say that Bitcoin is secure as long as 51% of the nodes are honest. The chain with the majority of nodes working on it will grow the fastest, thus will be the accepted chain. The difficulty of the new block is determined by an average, in order to generate blocks at a steady pace. There is also a Bitcoin reward associated with mining (publishing) a block. [5]

2.1.4 Merkle Tree

In the architecture of the blockchain, only the Merkle root is stored in the block header. Nodes only keep the recent blocks in memory. For the older blocks, they keep only the block header in memory. This storage ensures the integrity of the blockchain, while decreasing the memory required to have the full blockchain history. Since the hash of a block is the hash of the block header, this strategy does not impact the integrity checks of the blockchain. The Merkle root is the top of the Merkle tree, and is a unique identifier of the full tree. A Merkle tree is a tree where the parent node is the hash of the child nodes. The tree is immutable because changing a single node would have an impact on the Merkle root.

2.2 Marketplaces

The best way to buy Bitcoin for the first time is through marketplaces. Marketplaces facilitate the exchange of traditional currency for Bitcoin. However, it's important to understand that the Bitcoin you purchase initially remains in the custody of the platform, rather than being sent directly to your personal wallet. To gain custody of your Bitcoin, you need to request a transfer to your own wallet, similar to how traditional banking transactions rely on the bank to process your request. Once you have Bitcoins in your wallet, you can transact on the network without needing a third party. Unless you are running a node, you will need to trust a 3rd party, whether it is a marketplace, or over the counter, to first acquire Bitcoin.

When your Bitcoin is held in the marketplace's custody, it cannot be track onchain. It blends with the plarform's holdings. Marketplaces manage many wallets, some public and some private, which enhances the privacy of deposits and withdrawals but contradicts Bitcoin's design for transparency and independence from third parties. This lack of visibility poses a challenge, as users have no proof of the marketplace's solvency to reimburse all clients. However, this issue is being addressed through the introduction of proof of solvency (or proof of reserve) mechanisms, which allow marketplaces to demonstrate their solvency. While this is a positive development, current proof of solvency methods have shortcomings and are not fully sufficient to prove solvency.

2.3 Zero Knowledge

Zero-knowledge proof is a cryptographic techniques that allow a prover to demonstrate knowledge of a fact without divulging any additional information. For instance, rather than revealing the solution to an equation, zero-knowledge proofs enable the prover to show that they know the solution without revealing it. In this context, a witness is the secret information or evidence that supports the validity of the statement being proven. The goal is to construct a proof that is both sound and complete, and is zero-knowledge [7].

- **Completeness**: If the statement is true, an honest verifier will be convinced by an honest prover.

- **Soundness**: If the statement is false, no dishonest prover can convince the honest verifier (except with some infinitesimal probability).

- **Zero-Knowledge**: If the statement is true, a verifier learns nothing other than the fact that the statement is true. [31]

2.3.1 Non interactive proofs

Zero knowledge proofs were originally designed as interactive, that is, multiple rounds of interaction between the prover and the verifier [19]. leading to what are called interactive zero-knowledge proofs. This interaction allows the prover to demonstrate knowledge of the solution without revealing any additional information. An alternative model was then proposed where the verifier and prover use a reference string that is shared during a trusted setup. Once we have the reference string, a single message is needed between the prover and the verifier. The elimination of multiple rounds of interaction simplifies the verification process and reduces the computational power required. Therefore, noninteractive zero-knowledge proofs offer enhanced efficiency and scalability, which will be needed later on. [11] [18]

2.3.2 SNARKS

One of the recent advances for non-interactive proofs is what is known as SNARK (non-interactive argument of knowledge). This means a proof that is:

- **Succinct**: The size of the proof is very small compared to the size of the witness.

- **Non-interactive**: There are no rounds of interactions between the prover and the verifier.

- **Argument**: It is secure only against provers with bounded computational resources. A dishonest prover with unlimited computational power could potentially prove a false statement.

- **Knowledge-sound**: A valid proof can only be generated if the prover knows the witness. [29]

Moreover, a SNARK can also be zero-knowledge, where the prover demonstrates knowledge without revealing any additional information about the witness. We call such proof a zk-SNARK. There are many techniques to construct a SNARK, but all of them follow the same guideline:

- Express your problem as an arithmetic circuit

- Transform the circuit in a polynomial

- Commit the polynomial

- Interaction between the Prover and the Verifier to show that the polynomial solves the the problem

We will show an example of a SNARK using the technique known as Groth16[20]. To be able to construct the SNARK, we need to transform the code we want to prove in a quadratic arithmetic program (QAP).

Arithmetic circuit

Let's say we want to prove $x^3 + x + 5 = 35$. The Prover as to convince the Verifier that he knows the solution whitout revelaing it to him.

The first step of transforming a problem in the QAP form is to express a problem in the form of an arithmetic circuit. An arithmetic circuit is a set of gates, each assigned a distinct set of inputs corresponding to the numbers to be processed in the operation. These gates are configured to execute arithmetic operations such as addition, subtraction, multiplication, or division. The outputs of the gate circuit represent the digits of the resulting computation.


Figure 2.2: General Arithmetic circuit for

R1CS

The first step is to express our circuit as a set of constraint where each constraint contains only one arithmetic operation.

$$s_1 = x * x$$

$$s_2 = s_1 * x$$

$$s_3 = s_2 * x$$

$$out = s_3 + 5$$

We can now convert this into an R1CS. An R1CS is a sequence of groups of three vectors (a, b, c), where the solution to the R1CS is a vector s, and s must satisfy the equation $s \cdot a * s \cdot b - s \cdot c = 0$, where $\cdot$ represents the dot product. The length of each vector is the length of the total number of variables for the system. This includes a variable *one* at the beginning, the variable *out* at the end, and all intermediate variables. The standard way of representing the three vectors (a, b, c) is to create a mapping with the variables. Here is the mapping we will use:

$[one, out, x, s_1, s_2, s_3]$

This gives us for the first gate:

$$a = [0, 0, 1, 0, 0, 0]$$

$$b = [0, 0, 1, 0, 0, 0]$$

$$c = [0, 0, 0, 1, 0, 0]$$

We are checking here that $x * x = s_1$. a and b both represents x, while c represents s_1.

Let's verify that our first gate satisfies the solution equation:

$$s \cdot a * s \cdot b - s \cdot c = 0$$

We know that the solution to the equation is $x = 3$, so given that our s vector would be:

$$s = [1, out, x, s_1, s_2, s_3]$$

$$s = [1, 35, 3, 9, 27, 30]$$

Now if we do the dot product:

$$s \cdot a = [1, 35, 3, 9, 27, 30] \cdot [0, 0, 1, 0, 0, 0] = 0 * 1 + 0 * 35 + 1 * 3 + 0 *$$

$$9 + 0 * 27 + 0 * 30 = 3 \quad s \cdot b = [1, 35, 3, 9, 27, 30] \cdot [0, 0, 1, 0, 0, 0] = 3$$

$$s \cdot c = [1, 35, 3, 9, 27, 30] \cdot [0, 0, 0, 1, 0, 0] = 9 \quad s \cdot a * s \cdot b - s \cdot c = 3 * 3 - 9 = 0$$

This shows that our solution vector satisfies the first gate.

Following these rules, our second gate for $s_1 * x = s_2$ is:

$$a = [0, 0, 0, 1, 0, 0]$$

$$b = [0, 0, 1, 0, 0, 0]$$

$$c = [0, 0, 0, 0, 1, 0]$$

The third gate for $s_2 + x = s_3$:

$$a = [0, 0, 1, 0, 1, 0] \ (s_2 + x)$$

$$b = [1, 0, 0, 0, 0, 0] \ (1)$$

$$c = [0, 0, 0, 0, 0, 1] \ (s_3)$$

Fourth gate for $s_3 + 5 = out$:

$$a = [5, 0, 0, 0, 0, 1]$$

$$b = [1, 0, 0, 0, 0, 0]$$

$$c = [0, 1, 0, 0, 0, 0]$$

The R1CS put together:

$$A = \begin{bmatrix} 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & 0 \\ 5 & 0 & 0 & 0 & 0 & 1 \end{bmatrix}$$

$$B = \begin{bmatrix} 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

$$C = \begin{bmatrix} 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 & 0 \end{bmatrix}$$

[30]

18

QAP

Next step is converting our R1CS into QAP form, which uses polynomials instead of dot products. We do this by using polynomials instead of dot product. In order to get a polynomial, we need to do a Lagrange interpolation on a set of points. For a polynomial of degree n, we need $n + 1$ set points for the interpolation to give the polynomial. We need a polynomial of degree n, where n is the number of row - 1. (Same thing as number of gate -1). Therefore we need a set of 4 points to have a polynomial. Doing the transpose of our matrices, we can get 18 (3x6) polynomials. For instance, the first column of A^T is our first polynomial $[0, 0, 0, 5]$. This gives us the set of points $[(1, 0), (2, 0), (3, 0), (4, 5)]$. The polynomial interpretation of this set of points is $f(x) = 535x^3 + 636x^2 + 116x + 636$ Doing the same for every column of A, B and C, we get the matrices:

$$
A_m = \begin{bmatrix}
636 & 116 & 636 & 535 \\
0 & 0 & 0 & 0 \\
8 & 416 & 5 & 213 \\
635 & 330 & 637 & 321 \\
4 & 634 & 324 & 320 \\
640 & 536 & 640 & 107
\end{bmatrix}
$$

$$B_m = \begin{bmatrix} 3 & 529 & 323 & 427 \\ 0 & 0 & 0 & 0 \\ 639 & 112 & 318 & 214 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

$$C_m = \begin{bmatrix} 0 & 0 & 0 & 0 \\ 640 & 536 & 640 & 107 \\ 0 & 0 & 0 & 0 \\ 4 & 423 & 322 & 534 \\ 635 & 330 & 637 & 321 \\ 4 & 634 & 324 & 320 \end{bmatrix}$$

The polynomial that we got from the first column of A appears in the first row of A_m. Each row is obtained in a similar way.

The point of doing this transformation, is that now the dot product is a series of additions and multiplications of polynomials, and the result will be a polynomial.

The resulting polynomial need to equal 0 at all the x coordinates we used previously. If it does, it means that the checks pass. To check correctness, we do not need to evaluate our polynomial at every point. We can instead divide our polynomial t by another polynomial z, and verify that the division leaves no remainder. We define Z as $(x - 1) * (x - 2) * (x - 3) * ...$, the simplest polynomial that is equal to zero at

all points that correspond to logic gates.

In this case, computing T:

$$T(x) = S \cdot Am * S \cdot Bm - S \cdot Cm\ T(x) = 139 * x^6 + 372 * x^5 + 275 * x^4 +$$

$$58 * x^3 + 147 * x^2 + 379 * x + 553$$

In this case, we use $Z = (x-1)(x-2)(x-3)(x-4)$ To show that T is divided by Z, it is sufficient to verify $T(x) = 0$ at $1, 2, 3$ and 4.

More formally, we have shown that there exist a Polynomial H such that $T(x) = H(x) \cdot Z(x)$ (i.e. if T is divided by Z, then it will be perfectly divisible and will leave no remainder)

In a zkSNARK, the prover demonstrates that a polynomial Z exactly divides another polynomial T using Polynomial Commitments, such as KZG, Bulletproofs, or FRI. The purpose of this step is to reveal only certain evaluations of the polynomial, without disclosing the entire polynomial itself.[14]

2.3.3   Polynomial Commitments

A polynomial commitment is a cryptographic technique that allows to commit to (or "lock") a polynomial's data in such a way that the data can later be revealed (or "unlocked") while ensuring integrity and privacy.

Polynomial commitments are binding, meaning that once a commitment is made, it is computationally infeasible to alter the original polynomial without detection. The commitments are generally constant size. While collisions are theoretically possible

due to the pigeonhole principle, finding a second polynomial that matches the same commitment is computationally infeasible.

Additionally, polynomial commitments can be hiding, which means the verifier is not exposed to any information about the polynomial. For example, we can create a commitment by hashing a polynomial with a collision-resistant hash function like SHA-256, and only the hash value is shared. To make it hiding, a random factor could be added to the polynomial data before hashing.

Instead of committing to an entire dataset (a set of values or a table), we can commit to a polynomial that represents or encodes this data. This is useful because polynomials can be compact representations of structured data.

The naïve approach to committing to a polynomial is simply revealing its coefficients directly. This approach fails because the information is not hidden, among other things.

A well-designed polynomial commitment scheme allows flexibility: it enables someone to either open the commitment to reveal the entire polynomial, or to open it at a specific evaluation point. This is a necessary parameter of the scheme, as it allows verification that someone has knowledge of a polynomial that evaluates at a certain point without revealing the polynomial itself.

For the proof to be complete, the polynomial must be evaluated at $d+1$ unique points, where d is the degree of the polynomial. However, to be succinct, we want to verify the least number of points possible.

The question is, how many points do we need to verify to have a high confidence in

the polynomial? In a cryptographic setting, q is a large prime. If q is 256-bits and the polynomial is degree 1000, then the probability of giving the right value at random is $1000/2^256$ (As we recall, a polynomial of degree 1000 has 1000 roots). Therefore, after checking just one random point, it is statistically probable that the polynomials will be the same if the point matches. [35]

There are many different polynomial schemes, such as KZG, Bulletproofs and FRI. In these schemes, the setup processes and the commitments vary significantly.

KZG commitments rely on a trusted setup to generate elliptic curve parameters[23]. The KZG setup is universal. Once it is done, it can be useable by anyone. There are many different versions available to the public; the Ethereum Foundation has implemented one of these versions.

Bulletproofs avoids using a trusted setup by using logarithm-based techniques[15]. However, while the KZG proofs are constant size, Bulletproofs are logarithmic.

FRI evaluates polynomials through proximity checks.[8] FRI is post-quantum since it relies only on hash function security. Other schemes rely on cryptographic assumptions, such as the hardness of the discrete logarithm problem, which quantum computers can efficiently solve.

KZG

KZG commitments work by evaluating the polynomial at a single secret point τ. As we just saw, evaluating a single point is statistically sufficient. This approach ensures that the commitment remains succinct.

To achieve this, a trusted setup generates a structured random string (SRS) containing powers of τ of a specific polynomial degree.

$$\langle g^{\tau^0}, g^{\tau^1}, g^{\tau^2}, g^{\tau^3}, \ldots, g^{\tau^d} \rangle \equiv \text{SRS}$$

where g is a generator. The prover uses these precomputed values to commit to $f(x)$ without knowing τ or $P(\tau)$.

$$K_P(\tau) = \text{Commit}(P(\tau))$$

$$= g^{c_0}(g^{\tau})^{c_1}(g^{\tau^2})^{c_2}(g^{\tau^3})^{c_3} \ldots$$

$$= g^{c_0 + c_1\tau + c_2\tau^2 + c_3\tau^3 + \ldots}$$

$$= g^{P(\tau)}$$

BulletProofs

The commitment is succinct because the polynomial degree does not impact the size.

To open the polynomial, the prover can send the coefficients and the verifier will recompute $K_P(\tau)$. This is of little use since the prover could have simply hash the coefficient to get a similar result. KZG offers additionnal features that works specifically with polynomials.

Additive homomorphism in KZG allows commitments of polynomials to be com-

bined in a way that reflects polynomial addition. If two polynomials f and g have a sum $f + g = h$, then their commitments C_f and C_g have the sum $C_f + C_g = C_h$.

Multiplication is not exactly homormorphic in KZG, but we can get something close to it using bilinear pairing. The idea is to assert the value for $K_{P_1(\tau)} \cdot K_{P_2(\tau)}$ and convince it is correct given $K_{P_1(\tau)}$ and $K_{P_2(\tau)}$.

$$e(K_{P_1(\tau)}, K_{P_2(\tau)}) \stackrel{?}{=} e(K_{P_1(\tau) + P_2(\tau)}, g)$$

$$e(g^{P_1(\tau)}, g^{P_2(\tau)}) = e(g^{P_1(\tau) \cdot P_2(\tau)}, g)$$

$$= e(g, g)^{P_1(\tau) \cdot P_2(\tau)}$$

The most beneficial property of KZG, and any polynomial commitment scheme, is being able to open the commiment at specific points.

KZG commitments allow proving that a polynomial $P(x)$ has a root at r by demonstrating that $(x - r)$ divides $P(x)$ without remainder. As we recall, we can write any polynomial in a factorized format:

$P(x) = (x - r_0)(x - r_1)(x - r_2)\ldots$

If r is a root, $P(x)$ is divisible by $(x - r)$, so we define:

$Q(x) = \frac{P(x)}{x - r}$

The prover commits to $K_{Q(\tau)}$ and provides it to the verifier.

The verifier can compute $K_{V(\tau)}$ by treating $x - r$ as a polynomial:

$V(x) = x - r = -r + 1 * x$ and using KZG to produce it.

$$e(K_{P(\tau)}, g) \stackrel{?}{=} e(K_{Q(\tau)}, K_{V(\tau)})$$

$$e(g^{P(\tau)}, g) = e(g^{Q(\tau)}, g^{V(\tau)})$$

$$e(g, g)^{P(\tau)} = e(g, g)^{Q(\tau) \cdot V(\tau)}$$

This proof is independent of the polynomial degree of Q and P, making it efficient even for large polynomials.

The security of the commitments relies on properly generating the trusted setup. If the secret τ is leaked or manipulated, the security of the scheme could be compromised. To mitigate this, multi-party computation ceremonies ensure that at least one participant securely destroys their contribution to the SRS. Zero-knowledge proofs can also be used to validate that the setup was generated correctly without revealing τ.

KZG is the standard in terms of commitments for SNARKS. It is efficient and forms the basis of many cryptographic protocols. [34]

2.4 Proof of solvency

A proof of solvency is a system where we prove that an entity, in most cases an exchange, holds enough assets to cover the balance of every customer. In traditional finance, this would be done through an audit. While an audit can be useful, it has many limitations. Obviously it requires the trust of another 3rd party, but it also

poses a time constraint. Thus, it is impractical, even impossible to hold an audit every single day, and in the Bitcoin world things move fast. It is essential to be able to fill a proof of solvency every single day. The implementation of a mechanism to produce daily proofs of solvency is therefore of the utmost importance.

A proof of solvency is composed of a proof of assets, where we verify what assets the marketplace as control over, and the proof of liabilities, where we confirm that the total amount of user deposits is smaller than the marketplaces assets.

The first paper about a proof of solvency focuses only on the proof of liabilities. In his paper Gregory Maxwell addresses the issue of verifying the solvency of Bitcoin exchanges. [26] Maxwell's system ensures user privacy by maintaining confidentiality of individual account balances, while only revealing aggregate information in the proof of liabilities. This is achieved through the application of Merkle trees.


Figure 2.3: Merkle tree for proof of liabilities

In the Merkle tree, every node contains a user's balance along with a hash-based

commitment incorporating the customer ID and a nonce. The root of the tree is the sum of the balances. To verify their inclusion in the total liabilities, users receive a subset of the hash tree from the exchange. This subset includes the user's nonce and the sibling nodes along the unique path from the user's leaf node to the root. By comparing the received information to the exchange's broadcasted root node, users can confirm the inclusion of their balance. While elegant, the protocol does have privacy implications. The exact value of the exchange's total liabilities, published in the root node, may be sensitive data. Additionally, the proof of inclusion reveals the neighbor's balance, and the subtree's balance along the Merkle path.

The proof of liabilities is only part of the proof of solvency. A complete proof of solvency needs a proof of asset as well. Provision describes the first preserving privacy proof of asset [17].

In Provisions, the focus shifts towards preserving privacy while still proving ownership of assets. Instead of publicly demonstrating control over specific addresses, Provisions enables exchanges to prove ownership of an anonymous subset of addresses sourced from the blockchain.

Since these 2 papers, a lot of work was done to make the proof of solvency evolve. However, marketplaces still do not implement a proof of solvency, or a flawed and limited version of it.

2.4.1 Real world proof of solvency

In recent years, several major cryptocurrency exchanges, including Binance, Crypto.com, and Kraken, have taken steps to enhance transparency by provinding various proof of reserves. Binance has implemented a proof of reserves system where they publish a monthly Merkle tree as their proof of liabilities, and disclose a list of their assets (every cryptocurrencies under their control) [9]. Crypto.com published a one-time audit [27]. Kraken is also publishing a proof of liabilities every few months, without a proof of assets. [25].

Although these proof of reserves may seem promising initially, they are primarily superficial. The proofs have many shortcomings, they are not sufficient to prove that the marketplaces are solvent. The first concern is the lack of proof of assets, or in Binance's case the lack of evidence demonstrating control over the wallets. Without a reliable proof of assets, the proof of liabilities is worthless because it has nothing to compare against.

Moreover, the frequency of reporting is another area of concern. Given the dynamic nature of cryptocurrencies, a monthly report is not sufficient at all. Binance is the only marketplace describing their proof of solvency, and we can see that it is not built with recurrence in mind. The proof is created from scratch every time. They would need 150 servers to produce a daily proof [10].

Chapter 3

Recursion proofs

This chapter serves as an intermediary background exploration, diving into more advanced concepts beyond last chapter concepts. The realm of zero knowledge is dynamic and continuously evolving, marked by ongoing advancements and active research endeavors. Among these advancements, recursion proofs is one of the most important and active subjects. Recursion proofs are created to accelerate the generation of multiple zero-knowledge proofs. Recursion proofs differ in their design, they can vary significantly and serve different use cases. In this section, we will examine these variations, distinguishing between them. We aim to identify the most suitable method for our daily proof of liabilities and proof of inclusion. [22]

The prover's process in SNARKs is to create a proof using the setup parameters, a private witness, and public input. As discussed in the previous section, the proof begins by transforming the code into an arithmetic circuit, which is then represented as a polynomial. The polynomial encodes a trace of the circuit (every wire value

from inputs to intermediary values to output). This polynomial is created using the witness (private input).

The objective is to demonstrate that the polynomial satisfies the solution without revealing the polynomial itself. This is where the polynomial commitment scheme becomes essential. The setup phase generates public parameters that assist in creating and verifying the proof without disclosing the polynomial.

Table 3.1: Definitions of Symbols in SNARK Circuit

| Symbol	Definition
$S(C)$	Setup phase that generates the public parameters (pp and vp) for the prover and verifier.
$P(\text{pp}, x, w)$	Prover function that generates the proof π using the public parameters, public input x, and private witness w.
$V(\text{vp}, x, \pi)$	Verifier function that uses the verification parameters vp, public input x, and proof π to decide whether to accept or reject.
$C(x, w)$	Arithmetic circuit to generate the polynomial, taking public input x and private witness w.
w	Private witness, representing the prover's secret input.
x	Public input, accessible to both the prover and verifier.
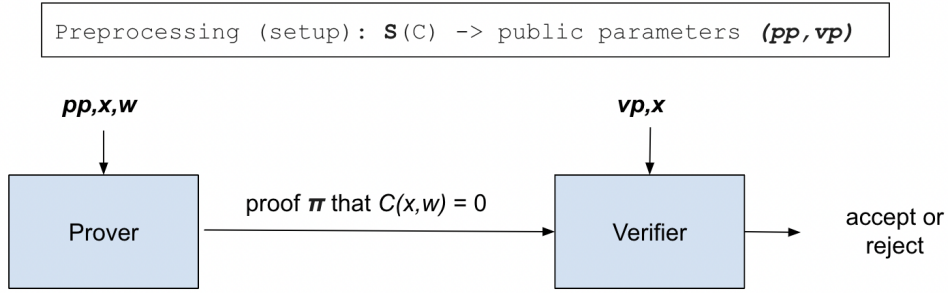
Figure 3.1: SNARK Circuit [1]

3.1 Aggregation

Aggregation is the simplest form of recursion for zk proofs. Its purpose is to reduce the number of proofs that must be verified. It takes a list of proofs as input and outputs a single proof. The process involves two phases.

In the first phase, you create your initial proofs. In the second phase, these proofs are combined into "proof of proofs". The second phase can be viewed as a tree, with the initial proofs being the leaves. Each parent node proves the proof of the child nodes. The tree can have any number of degrees. You can go from a tree of 2 levels to a binary tree with as many levels as needed.

This is a simple process. Each node has its proof, and then you prove that the other proofs are valid, giving you only one proof to verify. Since there are two different phases, two different circuits need to be constructed. On the surface, you can parallelize the initial proofs, which decreases the proving time, and you only have one proof to verify, which reduces the verifying time.

However, there are still some aggregation issues. The main issue is that a verification circuit is more complex than a proof of transaction. Verifying a proof involves constructing a circuit to represent the verification process itself, which is inherently complex. This circuit must perform cryptographic operations, including curve arithmetic and bilinear pairings, which must be implemented directly at the circuit level. As a result, the verification circuit is essentially the exponent of the cryptographic process. This becomes a scaling issue since the proof time of the second circuit grows linearly with the number of proofs to verify. [22]

3.2   Recursion scheme

The following technique is the recursion scheme, or Incrementally Verifiable Computation (IVC). Like the aggregation, the proof is separated into nodes. However, in this case each node serves two purposes. Provings that the node itself is valid, and that the previous proof is valid. The number of constraints will always stay the same because it only proves two things of fixed sizes. This creates a chain structure, where verifying the latest node is the equivalent of verifying every node in the chain. It is an improvement from the previous technique for two reasons. First, because this allows to maintain a constant number of constraints in the circuit. Second, because you don't have to redo a trusted setup once the size of the proof is defined. This means the trusted setup remains the same, and you don't have to redo it. This implies that the latest node does not take longer to verify than the second node. However, each

node still has to verify the previous node, which takes additional time. [22]

3.3   Accumulation

Accumulation schemes offer a different approach to aggregating proofs than IVC. Instead of verifying proofs at each step, accumulation defers all heavy computational tasks to the final step. This allows proofs to be efficiently added to an accumulator while minimalizing the intermediate verification workload. Each step adds a new proof to the accumulator while maintaining a small and constant-sized state. Unlike recursion, where each proof is verified within the circuit immediately, accumulation schemes delay all verification until the final stage, where all accumulated proofs are checked at once.

There are different types of accumulation schemes, each handling proof accumulation in a distinct way. Some, like polynomial accumulation, track commitments to polynomial evaluations over multiple steps. Others, such as accumulation based on cryptographic accumulators (e.g., RSA or Merkle accumulators), focus on aggregating statements into a compact structure. Regardless of the method, the core idea remains the same: verification is deferred to the final proof, preventing the verification cost from increasing with each step.

Accumulation schemes relate to Incrementally Verifiable Computation (IVC) because both aim to make verifying long computations more efficient. However, the key difference is that IVC requires each step to verify the previous one while keeping

proof sizes constant, while accumulation postpones all verification until the final step. This makes accumulation more suitable for cases where verifying everything at the end is preferable to verifying at every step.

The advantage is clear: instead of paying the cost of verification n times throughout the process, we pay it once at the end. This is especially useful when multiple independent proofs need to be efficiently combined without introducing the recursive overhead of IVC. However, this benefit comes with a tradeoff. Since all verifications are deferred, the final step can be computationally expensive.[16]

3.3.1 Halo accumulation

The concept of deferring the polynomial commitment opening, and consolidating them into a single operation, was introduced by Bowe, Grigg, and Hopwood Halo.[13] This process involves two parts: the first part is fast, where you output a polynomial and its commitment. The second part is expensive, where we verify the pair (f_1, c_1) of the polynomial f_1 and its commitment c_1. We open the polynomial only at a specific point. The second part can be accumulated if the commitment scheme is additively homomorphic. Instead of individually verifying each pair, we accumulate the pairs and verify their linear combinations ($c_1 + c_2 + ...$ is a commitment for $f_1 + f_2 + ...$). [35]

3.4   Folding scheme

In the Nova scheme, the folding technique accumulates the R1CS progressively as the computation proceeds rather than accumulating only the complex portions of the R1CS. This means that at each step in the folding process, instead of storing the separate R1CS sets for each node and combining them later, we incrementally combine them into a single R1CS set. This set, called "relaxed R1CS," simplifies the final proof generation. As a result, we avoid the need to handle multiple sets of R1CS in the final step, leading to a more efficient recursive proof process. The relaxed R1CS is then used to compute a single proof at the end of the folding process. [22] [24] [33]

3.4.1   Relaxed R1CS

Going back to the previous section, we previously defined what an R1CS is. The goal is to combine 2 R1CS and obtain another R1CS. If we are able to do that, we can fold every R1CS together and be left with only one.

If we **define our R1CS**:

fix an R1CS program $A, B, D \in \mathbb{F}_p^{u \times v}$

We define x as the public input, and w as the witness.

Instance 1: $x_1 \in \mathbb{F}_p^n$, $z_1 = (x_1, w_1) \in \mathbb{F}_p^v$

Instance 2: $x_2 \in \mathbb{F}_p^n$, $z_2 = (x_2, w_2) \in \mathbb{F}_p^v$

Where u is the number of constraints in the R1CS, v is the number of variables in the constraint system, p as the private field, and n the number

36

of public variables.

We know $Az_i \circ Bz_i = Dz_i$ for $i = 1, 2$

Recall: Our R1CS needs to be valid for any element of the field. To prove this, the verifier choose a random point r at which we will evaluate our constraints system.

The naïve approach to folding an R1CS is to sum the two instances together.

First attempt:

Lets define r as a random variable:

$r \leftarrow \mathbb{F}_p$

Then we set

$x \leftarrow x_1 + rx_2$

$z \leftarrow z_1 + rz_2 = (x_1 + rx_2, w_1 + rw_2)$

Then:

$$Az \circ Bz = A(z_1 + rz_2) \circ B(z_1 + rz_2)$$

$$= (Az_1) \circ (Bz_1) + r^2(Az_2) \circ (Bz_2) + (r(Az_2) \circ (Bz_1) + r(Az_1) \circ (Bz_2))$$

$$= Dz_1 + r^2 Dz_2 + E$$

Where E is a combination of the remaining values.

Simply summing the constraints does not preserve the R1CS structure we just defined: $Az_i \circ Bz_i = Dz_i$. Since the direct addition of two R1CS instances does not produce another valid R1CS, we need a different approach.

We need to modify the R1CS so that it can be folded. To solve this, we introduce a relaxed R1CS, which includes a new error term E. The goal is to have the new terms produced by the folding comprised in the term E. We will try the new form:
$Az_i \circ Bz_i = Dz_i E z_i$.

Let's **define a relaxed R1CS**:

$A, B, D \in \mathbb{F}_p^{u \times v}, (x \in \mathbb{F}_p^n, c \in \mathbb{F}_p, E \in \mathbb{F}_p^u)$

Witness: $z = (x, w) \in \mathbb{F}_p^v s.t. (Az) \circ (Bz) = c(Dz) + E$

Lets **fix the R1CS** program once again:

$A, B, D \in \mathbb{F}_p^{u \times v}$

Instance 1: public (x_1, c_1, E_1), witness $z_1 = (x_1, w_1) \in \mathbb{F}_p^v$

Instance 2: public (x_2, c_2, E_2), witness $z_2 = (x_2, w_2) \in \mathbb{F}_p^v$

We know $(Az_i) \circ (Bz_i) = c_i(Dz_i) + E_i$ for $i = 1, 2$

Second attempt:

$T \leftarrow (Az_2) \circ (Bz_1) + (Az_1) \circ (Bz_2) - c_1(Dz_2) - c_2(Dz_1)$

$x \leftarrow x_1 + rx_2, c \leftarrow c_1 + rc_2, E \leftarrow E_1 + rT + r^2 E_2$

$z \leftarrow z_1 + rz_2 = (x_1 + rx_2, w_1 + rw_2)$

$Az \circ Bz =$

$= A(z_1) \circ rB(z_1) + r^2(Az_2) \circ (Bz_2) + r(Az_2) \circ (Bz_1) + r(Az_1) \circ (Bz_2)$

$= c_1(Dz_1) + E_1 + r^2 c_2(Dz_2) + r^2 E_2 + r((Az_2) \circ (Bz_1) + (Az_1) \circ (Bz_2))$

38

$$= (c_1 + rc_2)(Dz_1 + rDz_2) + E_1 + r^2 E_2 + rT$$

$$= c(Dz) + E$$

We have a valid relaxed R1CS.

Validity While the error term E allows folding, it questions the validity of the proof. As long as E is bounded, the proof remains valid. The relaxed formulation preserves constraint structure, enables incremental folding, and prevents uncontrolled error propagation. The system can return to a strict R1CS if needed by proving $E = 0$ in the final verification, ensuring flexibility and correctness. [22]

Chapter 4

Proof construction

We will be using Circom to generate an arithmetic circuit, and SnarkJS for the proof generation.

Circom is a domain-specific language designed for creating arithmetic circuits, specifically utilized in zk-SNARKs. Within Circom, circuit code is written to define the desired constraints. One notable distinction from other languages is the utilization of signals and templates. Templates can be conceptualized as functions that operate as circuits. The primary template receives signals as inputs and produces other signals as outputs. Signals can be classified as either public or private. Assigning a value to a signal contributes to the constraint system and the witness calculation process. Input and outputs are signals, and you can have additionnal intermediate signals. During circuit compilation, constraints are generated in r1cs format. Additionally, compiling a circuit produces a witness file containing the data

essential for verifying the constraints and demonstrating the correct behavior of the circuit.

SnarkJS is a JavaScript library that provides tools for working with zk-SNARKs, including circuit compilation. SnarkJS facilitates generating zk-SNARK proofs for specific instances of the circuit using the r1cs and witness file. SnarkJS also provides utilities for verifying the proofs.

4.1 Proof of liabilities

The proof of liabilities operates on a list of balances and a list of email hashes as private inputs. The first purpose of the circuit is to validate that all values are non-negative and that all balances fall within a specified range. These verifications are crucial to prevent overflow or underflow issues, given that the operations occur within a finite field.

Subsequently, the proof of liabilities constructs a Merkle tree and outputs the total balance sum and the root hash of the Merkle tree. The pseudocode for the circuit is found here.

Inputs

- List of balance (private)

- List of email hash (private)

Outputs

- Balance Sum

- Root hash

- No negative values - boolean

- All small range - boolean

This proof of liabilities operates as intended because it returns the sum of the liabilities, which is exact because of the verifications. It also returns the root hash, ensuring you cannot alter any values inside the merkle tree. The merkle tree is hidden so that we do not give any information about users and their balances. The root hash will be used to verify the inclusion of the balances.

In a complete proof of reserves, the balance sum would be a private output. We would have another circuit proving that the sum of liabilities is smaller than the sum of assets, without revealing the balance.

4.2   Proof of inclusion

The proof of inclusion aims to prove that the balance of a user is included in the Merkle tree created in the proof of liabilities. To prove that a balance is included, it is sufficient to show that you know the Merkle path of a user balance, which we define using the list of neighbors sum, hash and binary. The neighbors binary variable indicates whether the neighbor is on the left or the right. The root hash, root sum,

user balance and user email hash are public because it needs to be shown which user is in which tree.

For the next figure, the blue nodes represent the value of the merkle path. We would have $neighborsSum = [29, 61]$, $neighborsHash = [Hash(User1), Hash(L2, R2)]$ and $neighborsBinary = [0, 1]$.
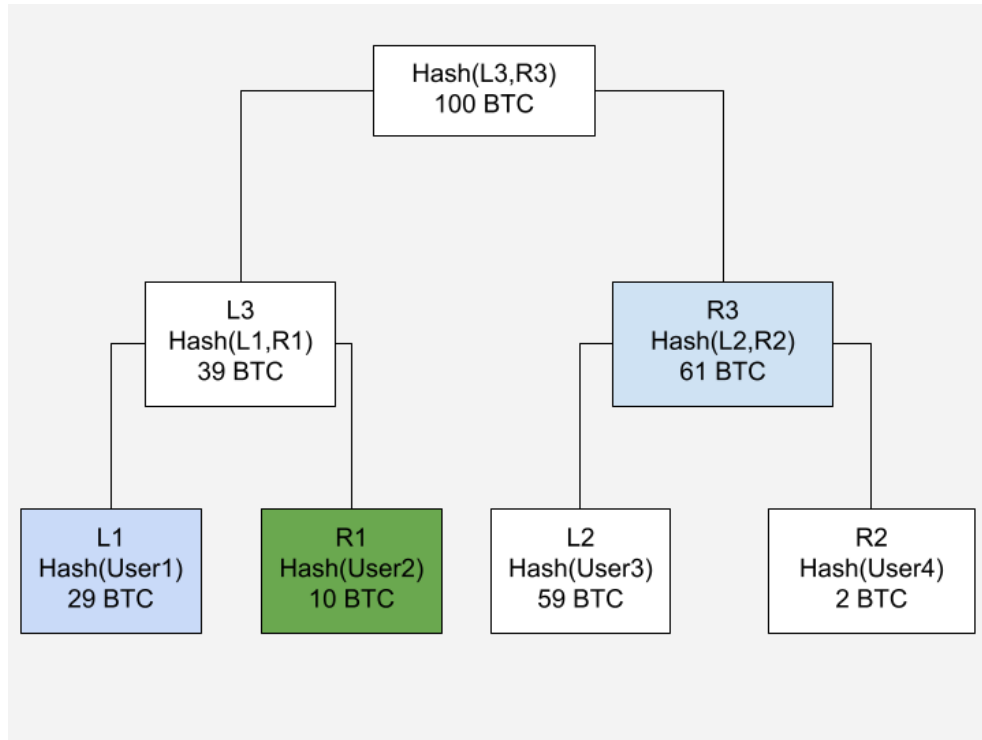

Figure 4.1: Merkle Path

Inputs

- List of neighbors sum (private)

- List of neighbors hash (private)

- List of neighbors binary (private)

- Root hash (public)

- Root sum (public)

- User balance (public)

- User email hash (public)

Outputs

- Balance included - boolean

In the circuit, we verify that the combination of the user balance, sum and merkle path gives the right root hash and root sum. There is no additional verifications since they are already done for this root hash, in the proof of liabilities.

4.3   Daily proof of liabilities

We aim to enhance our current circuit by minimizing the work needed in subsequent rounds. The first thing to explore is recursion proofs, as they were specifically created to reduce the total computational effort across rounds.

4.3.1   Aggregation proofs

The main advantage of the aggregation proof is in streamlining the verification process. For instance, in the second round, we can prove the integrity of the current round and all previous rounds. However, this benefit comes with trade-offs. The first

drawback is the increase in proof size, as it necessitates to prove the current circuit and verify the previous ones. When considering the frequency of verification, having a fixed number of nodes verify the proof daily, it would be illogical to make the nodes verify the previous rounds every single day. On the other hand, if the verification is not consistent, or there is a need for new nodes to be able to come in and verify every proof quickly, then aggregation becomes more appealing.

In our case, the priority is to produce daily succinct proofs. We need to ensure the integrity of every round, while keeping the proof size to a minimum. Therefore, having our nodes verify the circuit at every round without the computational overhead of the aggregated proof is sufficient.

4.3.2   Other recursion proofs

The aggregation proof stands out as the only recursion proof potentially useful in certain scenarios. If we examine the other types of recursion proofs, namely Recursion scheme, accumulation and Folding scheme, they all have one thing in common. They are all designed to verify multiple rounds concurrently. This approach is not aligned with our objectives. We are interested in working on rounds independently and reducing the individual workload.

4.3.3   Change circuit

If we cannot use any recursion schemes, we need an alternative approach to reduce the complexity of subsequent rounds. Our solution is to reutilize the same Merkle

tree as the previous rounds, modify and adapt it to include the changes. The key challenge is that the Merkle tree was built inside the circuit, and is therefore not accessible.

What we will be doing, in our modified circuit, is adjust the Merkle tree inside the circuit. For every change, we send the corresponding Merkle Path which will be verified by the circuit. The circuit will then compute a new Root Hash for each change and output the final Markle Hash.

The standard verification will be applied to the new values (e.g., non-negative values, limited range).

Inputs

- List of old email hash (private) - 1 per change

- List of old values (private) - 1 per change

- List of new email hash (private) - 1 per change

- List of new values (private) - 1 per change

- List of temporary root hash (private) - 1 per change

- List of temporary root sum (private) - 1 per change

- Old root hash (public)

- Old root sum (public)

- List of neighbors sum (private) - 1 list per change

- List of neighbors hash (private) - 1 list per change

- List of neighbors binary (private) - 1 list per change

Outputs

- Valid hash

- Valid sum

- No negative values

- All small range

- New root hash

- New root sum

Here is an example of a change. In this case the change is a change in balance. First we prove that the 10 BTC of user 1 are included in the merkle tree. Then we prove that the 11 BTC of user 1 are included in the new merkle tree. After the last change we are left with the new root balance and the new root hash. Not that for every change we only need 1 merkle path.

For every change in the merkle tree, we have the merkle path with the old values, the new values, the temporary root hash and the temporary root sum. The circuit is iterating over the changes and gives a final root hash and final root sum.

(a) Pre change   (b) Post change

Figure 4.2: Merkle tree change

Minimizing the number of changes

To minimize the number of constraints, we want to minimize the number of changes. We define a change as a hash change at the leaf level. Any new user, balance change or removal of old user is considered as a change. The naive change calculation would be $changes = balanceChanges + newUsers + deprecateUser$. However, a new user can take the node of a deprecater user. So we would have the equation $changes = balanceChanges + max(newUsers + deprecateUser)$.

Constraint analysis

Theorically, the new circuit makes sense. It should be quite faster than the original one, so let's analyze it. We want to compare the number of constaints for the original circuit with the number of constraints with the change circuit. The first assumption we are gonna need to do, is to assume the number of changes in a day. A completely arbitrary value would be 1%.

Figure 4.3: Number of constraints original circuit vs new circuit with 1% change

Once you reach more than 128 users with 1% change, it is not worth it to use the new circuit. This is not really practical because any marketplace will have more than 128 users. However, we can see that the number of constraints of the new circuit grows linearly for the number of change. This means that two circuits of 0.5% change would have a similar number of constraints as the one circuit with 1% change. We can take this a step further, and instead of looking at daily proof look at hourly proof. Going with the assumption that 1% change daily, we will assume that 0.05% of the balances changes hourly. Lets look at the number of constraints if 0.05% of users change every hour.

Figure 4.4: Number of constraints (millions) original circuit vs new circuit with 0.05% change

Even with more than 1 million users, there is still 3 times less constraints with our new circuit. So if we want to do an hourly proof, it is less expensive with the new circuit than the old circuit. We can take this a step further, and do a proof at every new block. We could have a merkle tree up to date for the liabilities side at every single block.

While the hourly proof and block proof is better with the new circuit, maybe a daily proof is sufficient for most use cases. It is still less work to do the daily proof with the original circuit, than to do 24 hourly proof with the new circuit.

4.4 Daily proof of inclusion

Because a marketplace can have millions of users, it is impractical to build a proof of inclusion for every single user on a daily basis. It is the user's responsibility to request a proof that their balance is included in the published Merkle tree.

In an ideal world, a Merkle tree would be published at least daily. However, it is once again impractical to require every user to verify their inclusion everyday. Nevertheless, each user verifying its inclusion in the tree increases the chance that the proof of liabilities is valid and includes every single balance. This is why it is primordial to find a way to make it easier to verify the inclusion. This is where Nova folding schemes come in.

The novel way to do proof of inclusion is to generate a proof of inclusion that is valid from the day of creation of the account, to the requested day by the user. Normally you would need to create 500 different proofs for 500 days. However, we saw in the previous section that the nova folding scheme enables the 500 proofs into just one. Verifying this one proof is the equivalent of verifying the 500, or any other number of days required. This drastically simplifies the verification process.

Previously, when requesting a proof of inclusion, the user would only receive proof that their balance is included in the latest published Merkle tree. Now, the proof verifies that the balance has been included in every previously published Merkle tree. This ensures that any malicious entity would be unable to alter ownership of balances across multiple days. It might seem like a small detail, but it is the detail that makes

all the difference, and here is why.

In this chart, we can observe the failure probability, which represents the likelihood that a dishonest prover gets away with misbehaving. If we take for example the orange line, where an exchange has 1 million users, we notice that the failure probability approaches 0 when over 4% of users perform their verifications. However, without utilizing folding, this principle doesn't extend to every round. It implies that for each round, we require a minimum of 4% of users to have a proof with significant confidence.

(c) $\tau = 0, c = 0.01\% \cdot n$ and varying n

Figure 4.5: Failure Probability [21]

Now, lets evaluate the failure probability if we are using folding. In the initial round, with no changes, the failure probability for 20k users remains at 13%.

Figure 4.6: Failure Probability Round 1 [12]

However, if we have 20k distinct users in round 2, the failure probability of round 1 decreases to 1.6% .

Figure 4.7: Failure Probability Round 1 After 2 Rounds[12]

If we have 20k different users in round 3, the failure probability of round 1 further

decreases to 0.2% .

Figure 4.8: Failure Probability Round 1 after 3 Rounds [12]

The key takeaway here is that without folding, you require 4% of users to verify every round. However, with folding you only need 4% of users to participate in at least 1 round. This significantly reduces the burden on the user to verify, while increasing the confidence in the proof.

4.4.1 Circuit Design

The private inputs vary for each instance, whereas the public inputs are carried over from round to round. In order to implement folding, we need to slightly adjust our circuit. Everything except the way we handle inputs and outputs stays the same. The private inputs vary for every instance, while the public inputs are carried over from round to round.

Inputs

- List of neighbors sum (private)

- List of neighbors hash (private)

- List of neighbors binary (private)

- Root hash (private)

- Root sum (private)

- User balance (private)

- User email hash (private)

- Steps in (public)

Outputs

- Steps out (public)

Steps

- Balance included

- Root sum

- Root hash

- User balance

- User email hash

All the data that is passed around between the circuits is public and is called step in and step out, and the regular inputs are private. The step in of a circuit is the step out of the previous one.


Figure 4.9: Folding Circuit

S0 is the step in of the first change, S1 is the step out of first change and step in of second change. F is the folded circuit and T is the number of time it is folded.

These steps encompass all the public values we initially had in the circuit. None of the variables of the steps are used in the circuit itself. They are used as a means to make a value public.

The initial circuit takes meaningless values as inputs, while the following circuits

take the public values of the previous circuit. At the end you only have to verify a single proof. You also need to compare the circuit output with the proof of liabilities outputs for every round.

4.5   Folded daily proof of liabilities

With the original circuit, we saw how it was not usefull to use folding or any other recursion scheme. However, with the new circuit we saw that we were able to separate a big proof into multiple smaller ones. This aligns prefectly wth the recursions ideas. We can separate our new circuit of n changes into m circuits, where $m <= n$, and fold the circuits together.


Figure 4.10: Proof time original circuit vs folded circuit with 1% change, 1 change by fold

As we can see, after 10 000 users, both circuits seems to be groing linearly, with the folded circuit doing much better. For instance, at 1M users, we have the old circuit taking 196608 seconds, and the new circuits taking 17188 seconds. This is more than 10 times better. We should take the number of seconds with a grain of salt, since these measurements were done on my mac, and there are much more performant computers out there. The data we should focus on is that it took 10x less time for the folded circuit, and it was not event optimized. We were doing 1 change by circuit, so at 1M users we folded 10k circuits.

Now that we know that our circuit is better, we can optimize it.


Figure 4.11: Optimization of the proof time and verifier time of the folded circuit with 1% change and 1M users

As we can see, the optimize time for the proofer, is between 8 and 32 changes by fold, while the verifier time seems to be groing at a monotonic pace. Depending on

the use case, you can decide the numbe of changes you should have for every fold.

4.5.1   Circuit design

The folded change circuit is exactly the same as the regular change circuit, but the inputs works differently because some data is passed around the circuits. There is 2 types of step variables; the variables needed in the following circuit, and the variables that are just there to be public. Starting from the inputs of the regular change circuit, the old hash and old sum root are moved to the step in, while all the outputs are moved to step out.

Inputs

- List of old email hash (private) - 1 per change

- List of old values (private) - 1 per change

- List of new email hash (private) - 1 per change

- List of new values (private) - 1 per change

- List of temporary root hash (private) - 1 per change

- List of temporary root sum (private) - 1 per change

- List of neighbors sum (private) - 1 list per change

- List of neighbors hash (private) - 1 list per change

- List of neighbors binary (private) - 1 list per change

- Step in (public)

Outputs

- Steps out

Steps

- Valid hash and valid sum

- No negative value and all small range

- Root hash

- Root sum

Chapter 5

Conclusion

This thesis introduces a novel approach to generating a more compact proof of liabilities. Leveraging a fresh circuit design, along with leveraging prior proofs and the folding scheme, we achieved notable reductions in proof size. This smaller proof position itself well to be generated daily, and even at higher frequencies. Our analysis demonstrates the superior performance of the folding scheme combined with the updated circuit, compared to the existing model. The significance of a more concise proof of liabilities cannot be overstated, serving as a cornerstone for broader adoption of robust proof of solvency.

To complete the proof of solvency for marketplace integration, a proof of assets is imperative. One challenge in this endeavor lies in conducting proofs for multiple currencies. However, by employing the folding scheme, we can segment the proof of assets into currency-specific proofs before conducting the folded proof, mitigating this complexity.

Appendix A

Sample Code

This section contains the pseudocode for the circuits used in both the proof of liabilities, and proof of inclusion.

A.1 Proof of liabilities

Listing A.1: Liabilities circuit pseudocode

```
# Define a function for the Merkle tree sum circuit
def sum_merkle_tree(levels, inputs, balances[inputs],
                    email_hashes[inputs]):
    # Outputs variables for sum, root hash, and range checks
    root_sum = 0
    root_hash = 0
    not_negative = False
```

```python
all_small_range = False


# Define signals lists for sum and hash nodes at each level
sum_nodes = [[0] * inputs for _ in range(levels+1)]
hash_nodes = [[0] * inputs for _ in range(levels+1)]


# Initialize variables
level_size = inputs
max_bits = 100
temp_not_big = 0
temp_not_negative = 0


# Loop through each input
for i in range(inputs):
    # Assign input values to hash and sum nodes
    hash_nodes[0][i] = email_hashes[i]
    sum_nodes[0][i] = balances[i]


    # Perform range check and negative check
    rangecheck = RangeCheck(maxBits, balances[i])
    tempNotBig += rangecheck
    negativecheck = NegativeCheck(newValues[i])
```

```python
        tempNotNegative += negativecheck


    # Check if all balances are within a small range
    if temp_not_big == inputs:

        all_small_range = True


    # Check if all balances are non-negative
    if temp_not_negative == inputs:

        not_negative = True


    # Loop through each level
    for i in range(levels):
        # Loop through each pair of nodes at the current level
        for j in range(0, level_size, 2):
            # Store sum and root hash for the next level
            sum_nodes[i+1][nextLevelSize] =

            sum_nodes[i][j] + sum_nodes[i][j+1]

            hash_nodes[i+1][nextLevelSize] =

            hash_nodes[i][j] + hash_nodes[i][j+1]


        # Update the size of the current level
        levelSize = nextLevelSize;
```

```
        nextLevelSize = 0;


    # Assign final sum and root hash values

    root_sum = sum_nodes[levels][0]

    root_hash = hash_nodes[levels][0]


    return root_sum, root_hash, not_negative, all_small_range
```

A.2 Proof of inclusion

Listing A.2: Inclusion circuit pseudocode

```
# Define a function for the inclusion proof circuit

def inclusion(levels, neighborsSum, neighborsHash, neighborsBinary,

              step_in, sum, rootHash, userBalance, userEmailHash):
    # Initialize sum and hash nodes

    sumNodes = [0] * (levels+1)

    hashNodes = [0] * (levels+1)

    sumNodes[0] = userBalance

    hashNodes[0] = userEmailHash


    # Iterate through each level

    for i in range(levels):
```

```
    # Update hash node

    if neighborsBinary[i]:

        hashNodes[i+1]=getHash(neighborsHash[i],

        neighborsSum[i],hashNodes[i],sumNode[i])

    else:

        hashNodes[i+1]=getHash(hashNodes[i],sumNode[i],

        neighborsHash[i],neighborsSum[i])

    # Update sum node

    sumNodes[i + 1] = neighborsSum[i] + sumNodes[i]


# Check validity of root hash

validHash = IsEqual([hashNodes[levels], rootHash])


# Check validity of sum

validSum = IsEqual([sumNodes[levels], sum])


return validSum,validHash
```

A.3   Change circuit

Listing A.3: Liabilities change circuit pseudocode

```
# Define a function for the liabilities change proof circuit
```

```python
def liabilities(levels, changes, oldEmailHash[changes],
        oldValues[changes], newEmailHash[changes], newValues[changes],
        tempHash[changes+1], tempSum[changes+1],
        neighborsSum[changes][levels], neighborsHash[changes][levels],
        neighborsBinary[changes][levels]):
    # Calculate newRootHash and newSum
    newRootHash = tempHash[changes+1]

    newSum = tempSum[changes+1]

    oldSum = tempSum[0]

    oldRootHash = tempHash[0]

    currentSum = oldSum


    # Part 1: Check validity of new values
    tempNotBig = 0

    tempNotNegative = 0

    maxBits = 100


    # Iterate through each change
    for i in range(changes):
        # Calculate currentSum
        currentSum += newValues[i] - oldValues[i]
```

```
# Perform range check and negative check
rangecheck = RangeCheck(maxBits, newValues[i])
tempNotBig += rangecheck
negativecheck = NegativeCheck(newValues[i])
tempNotNegative += negativecheck


# Check if all new values are within range
allSmallRange = IsEqual([changes, tempNotBig])


# Check if all new values are not negative
notNegative = IsEqual([changes, tempNotNegative])


# Check if newSum equals currentSum
equalSum = IsEqual([newSum, currentSum])


# Part 2: Check validity of old and new paths
# Ensure that old root + change = temp root
tempValidHash = [0] * (changes + 1)
tempValidSum = [0] * (changes + 1)
tempOldHashEqual = [0] * (changes + 1)
tempOldSumEqual = [0] * (changes + 1)
tempValidHash[0] = 1
```

```python
        tempValidSum[0] = 1


#For each level, we are veryfing the inclusion of the value changing
#and the new value.
#hashNode[0] and sumNodes[0] is for the inclusion of the value changing
#hashNode[1] and sumNodes[1] is for the inclusion of the new value


    for j in range(changes):
        for i in range(levels):
        if neighborsBinary[j][i]:
            hashNodes[0][j][i+1]=getHash(neighborsHash[j][i],
            neighborsSum[j][i],hashNodes[0][j][i],sumNode[0][j][i])
            hashNodes[1][j][i+1]=getHash(neighborsHash[j][i],
            neighborsSum[j][i],hashNodes[1][j][i],sumNode[1][j][i])
        else:
            hashNodes[0][j][i+1]=getHash(hashNodes[0][j][i],
            sumNode[0][j][i],neighborsHash[j][i],neighborsSum[j][i])
            hashNodes[1][j][i+1]=getHash(hashNodes[1][j][i],
            sumNode[1][j][i],neighborsHash[j][i],neighborsSum[j][i])


        sumNodes[0][j][i+1] = sumNode[0][j][i] + neighborsSum[j][i]
        sumNodes[1][j][i+1] = sumNode[0][j][i] + neighborsSum[j][i]
```

```
# Old  calculated  old  hash  is  in  tempHash

hashEqual  =  IsEqual ( hashNodes [ 0 ] [ j ] [ levels ] ,  tempHash [ j ] )

tempOldHashEqual [ j+1]  =  tempOldHashEqual [ j ]  *  hashEqual


# New  temp  hash  is  valid

hashEqual  =  IsEqual ( hashNodes [ 1 ] [ j ] [ levels ] ,  tempHash [ j+1])

tempValidHash [ j+1]  =  tempValidHash [ j ]  *  hashEqual


# Old  calculated  sum  is  in  tempSum

sumEqual  =  IsEqual ( sumNodes [ 0 ] [ j ] [ levels ] ,  tempSum [ j ] )

tempOldSumEqual [ j+1]  =  tempOldSumEqual [ j ]  *  sumEqual


# New  temp  sum  is  valid

sumEqual  =  IsEqual ( sumNodes [ 1 ] [ j ] [ levels ] ,  tempSum [ j+1])

tempValidSum [ j+1]  =  tempValidSum [ j ]  *  sumEqual


validHash  =  tempValidHash [ changes ]  *  tempOldHashEqual [ changes ]

validSum  =  tempValidSum [ changes ]  *  tempOldSumEqual [ changes ]


return  ( notNegative , allSmallRange , validHash ,
```

validSum , newRootHash , newSum)

Bibliography

[1] Zero knowledge proofs. MOOC, 2023. Dan Boneh, Shafi Goldwasser, Dawn Song, Justin Thaler, Yupeng Zhang.

[2] Zero knowledge proofs. MOOC, 2023. Dan Boneh, Shafi Goldwasser, Dawn Song, Justin Thaler, Yupeng Zhang.

[3] Zero knowledge proofs. MOOC, 2023. Dan Boneh, Shafi Goldwasser, Dawn Song, Justin Thaler, Yupeng Zhang.

[4] L. A. An incomplete guide to folding: Nova, sangria, supernova, hypernova, protostar. https://taiko.mirror.xyz/tk8LoE-rC2wOMJ4wCWwaJwbq8-Ih8DXnLUf7aJX1FbU, August 2023.

[5] A. M. Antonopoulos. Mastering bitcoin: Unlocking digital cryptocurrencies. O'Reilly Media, 2017. 2nd edition.

[6] Bake. Merkle tree proof of reserves and liabilities: Raising the bar for enhanced transparency, Dec 2022.

[7] B. Barak. Lecture 14: Zero knowledge proofs. *Boaz Barak's Blog*.

[8] E. Ben-Sasson, I. Bentov, Y. Horesh, and M. Riabzev. Fast reed-solomon interactive oracle proofs of proximity. 2023. Technion, Haifa, Israel and Cornell University, Ithaca, NY, USA.

[9] Binance. Proof of reserves. *Binance*.

[10] Binance. Proof of solvency. *Binance*.

[11] M. Blum, P. Feldman, and S. Micali. Non-interactive zero-knowledge and its applications. page 103–112, 1988.

[12] E. Bottazzi. Zk10: Incremental proof of solvency with nova. Presentation at ZK10 Summit, September 30 2023. Video: https://www.youtube.com/watch?v=sRAA1RYYHEs.

[13] S. Bowe, J. Grigg, and D. Hopwood. Recursive proof composition without a trusted setup. *Electric Coin Company*, 2023.

[14] V. Buterin. Quadratic arithmetic programs: From zero to hero. *Medium*, December 2016.

[15] B. Bünz, J. Bootle, D. Boneh, A. Poelstra, P. Wuille, and G. Maxwell. Bulletproofs: Short proofs for confidential transactions and more. pages 319–334, 2018. Accessed: 2024-09-08.

[16] B. Bünz, A. Chiesa, P. Mishra, and N. Spooner. Proof-carrying data from accumulation schemes. *Cryptology ePrint Archive*, (2020/499), 2020. Accessed: 2025-01-29.

[17] G. G. Dagher, B. Bunz, J. Bonneau, J. Clark, and D. Boneh. Provisions: Privacy-preserving proofs of solvency for bitcoin exchanges. 2015.

[18] O. Goldreich, S. Micali, and A. Wigderson. Interactive and noninteractive zero knowledge are equivalent in the help model? page 113–131, 1991.

[19] S. Goldwasser, S. Micali, and C. Rackoff. The knowledge complexity of interactive proof systems. *SIAM Journal on Computing*, 18(1):186–208, 1989.

[20] J. Groth. On the size of pairing-based non-interactive zero-knowledge proofs. In *Proceedings of the 2016 ACM Conference on Computer and Communications Security (CCS)*, pages 480–491. ACM, 2016.

[21] Y. Ji and K. Chalkias. Generalized proof of liabilities, 2021.

[22] z. L. Jim.ZK. Nova studies i: Exploring aggregation, recursion, and folding. *zkLinkBlog*, Nov 2023.

[23] A. Kate and G. M. Zaverucha. Polynomial commitments. December 2010. Accessed: 2024-09-08.

[24] A. Kothapalli, S. Setty, and I. Tzialla. Nova: Recursive zero-knowledge arguments from folding schemes. 2021.

[25] Kraken. Proof of reserves.

[26] D. Malkhi. Exploring proof of solvency and liability verification systems. 2023.

[27] Mazars. Proof of reserve report. 2022.

[28] S. Nakamoto. Bitcoin: A peer-to-peer electronic cash system. *Bitcoin.org*, 2008. Technical report.

[29] A. Nitulescu. zk-snarks: A gentle introduction. 2020.

[30] RisenCrypto. R1cs and qap, 2023.

[31] L. Team. Introduction to zero-knowledge proofs. *LCX*, November 2023.

[32] Trace. The proof supply chain. *Figment Capital*, January 2024.

[33] E. van Oorschot, Y. Deng, and J. Clark. Folding (sangria). GitHub Pages, 2024.

[34] E. van Oorschot, Y. Deng, and J. Clark. Kzg commitments. GitHub Pages, 2024.

[35] Veridise. Halo and accumulation. *Veridise*, August 2023.

[36] Veridise. Nova and folding (1/2). *Veridise*, September 2023.

[37] Veridise. Recursive snarks and incrementally verifiable computation (ivc) part i: Recursive snarks and incrementally verifiable computation. *Veridise*, July 2023.